

# Housing\_Data\_Assignment

Theodore Koby-Hercsky

5/9/2021

html\_document: <https://rpubs.com/theoKoby/769360>

## Set the working directory to the root of your DSC 520 directory

```
setwd("~/Documents/Bellevue University Classes/DSC520/assignments/assignment06")
```

## The housing.csv file

*## I am importing readr from the library so I can use the read\_csv function to create my student survey data frame.*

```
library(readr)
```

*## Creating the student survey data frame by using the read\_csv function to pull my student survey data.*

```
Housing_df <- read_csv("data/housing.csv")
```

```
Housing_df
```

```
## # A tibble: 12,865 x 24
```

```
##   sale_date sale_price sale_reason sale_instrument sale_warning sitetype  
##   <chr>         <dbl>         <dbl>         <dbl> <chr>         <chr>
```

```
## 1 1/3/06         698000             1             3 No           R1
```

```
## 2 1/3/06         649990             1             3 No           R1
```

```
## 3 1/3/06         572500             1             3 No           R1
```

```
## 4 1/3/06         420000             1             3 No           R1
```

```
## 5 1/3/06         369900             1             3 Yes          R1
```

```
## 6 1/3/06         184667             1            15 Yes          R1
```

```
## 7 1/4/06        1050000             1             3 No           R1
```

```
## 8 1/4/06         875000             1             3 No           R1
```

```
## 9 1/4/06         660000             1             3 No           R1
```

```
## 10 1/4/06        650000             1             3 No           R1
```

```
## # ... with 12,855 more rows, and 18 more variables: addr_full <chr>,
```

```
## #   zip_code <dbl>, ctynome <chr>, postalctyn <chr>, lon <dbl>, lat <dbl>,
```

```
## #   building_grade <dbl>, square_feet_total_living <dbl>, bedrooms <dbl>,
```

```
## #   bath_full_count <dbl>, bath_half_count <dbl>, bath_3qtr_count <dbl>,
```

```
## #   year_built <dbl>, year_renovated <dbl>, current_zoning <chr>,
```

```
## #   sq_ft_lot <dbl>, prop_type <chr>, present_use <dbl>
```

```
#Explain any transformations or modifications you made to the dataset
```

I modified the sales warning variable by changing the numbers that are useless to us to being a yes for a sales warning and no if there was no data in that field.

I also inputted the city name if the line was blank in the city name variable by checking the name on the postal city name variable.

Create two variables; one that will contain the variables Sale Price and Square Foot of Lot (same variables used from previous assignment on simple regression) and one that will contain Sale Price and several additional predictors of your choice. Explain the basis for your additional predictor selections.

*## I fit a linear model using the `sq\_ft\_lot` variable as the predictor and `sale\_price` as the outcome*

```
square_foot_price_lm <- lm(sale_price ~ sq_ft_lot, data = Housing_df)
```

*## When we fit our linear model by using our sq\_ft\_lot variable as the predictor and sale\_price as the outcome with the Housing\_df as our data we see coefficients for intercept at 6.418*

```
square_foot_price_lm
```

```
##
```

```
## Call:
```

```
## lm(formula = sale_price ~ sq_ft_lot, data = Housing_df)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)    sq_ft_lot
```

```
##  6.418e+05    8.510e-01
```

*## Seen below is a linear model using the `zip\_code`, `square\_feet\_total\_living`, `bedrooms`, and `bath\_full\_count` variables as the predictors and `sale\_price` as the outcome*

```
sale_price_lm <- lm(sale_price ~ zip_code + square_feet_total_living + bedrooms + bath_full_count, data = Housing_df)
```

*## We fit the linear model by using our variables as the predictor and sale\_price as the outcome with the Housing\_df as our data we see coefficients for intercept and `zip\_code`, `square\_feet\_total\_living`, `bedrooms`, and `bath\_full\_count` which is the slope for the predictors.*

```
sale_price_lm
```

```
##
```

```
## Call:
```

```
## lm(formula = sale_price ~ zip_code + square_feet_total_living +
```

```
##   bedrooms + bath_full_count, data = Housing_df)
```

```
##
```

```
## Coefficients:
##              (Intercept)              zip_code  square_feet_total_livi
ng
##              -1.999e+08              2.041e+03              1.837e+
02
##              bedrooms              bath_full_count
##              -2.466e+04              4.195e+04
```

Execute a `summary()` function on two variables defined in the previous step to compare the model results. What are the R2 and Adjusted R2 statistics? Explain what these results tell you about the overall model. Did the inclusion of the additional predictors help explain any large variations found in Sale Price?

*## I will view the full report by using summary of my square\_foot\_price\_lm model*

*## As seen below we see that the Multiple R-squared: is 0.01435 and the Adjusted R-squared: is 0.01428*

*## Which is expressed as a percentage between 0 and 100, with 100 signaling perfect correlation and zero no correlation at all. As in the summary of the square\_foot\_price\_lm model we find that the Multiple R-squared and Adjusted R-squared are less than 2% which is very low meaning there is very little correlation. Which I find to be quite interesting as you would expect the square footage to have a higher impact on the price.*

```
summary(square_foot_price_lm)
```

```
##
## Call:
## lm(formula = sale_price ~ sq_ft_lot, data = Housing_df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2016064  -194842   -63293    91565   3735109
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  6.418e+05  3.800e+03  168.90  <2e-16 ***
## sq_ft_lot    8.510e-01  6.217e-02   13.69  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 401500 on 12863 degrees of freedom
## Multiple R-squared:  0.01435,    Adjusted R-squared:  0.01428
## F-statistic: 187.3 on 1 and 12863 DF,  p-value: < 2.2e-16
```

*## The Summary for sale\_price\_lm can be seen when we use the Summary function seen below*

*## As seen below we see that the Multiple R-squared: is 0.2121 and the Adjusted R-squared: is 0.2118*

*## As for the summary of the sale\_price\_lm model we find that the Multiple R-squared and Adjusted R-squared are significantly higher than that of the square\_foot\_price\_lm model*

*re\_foot\_price\_lm meaning there is a higher correlation in comparison to the previous test. While in the inclusion of the additional predictors helped explain the variance in sales price as the zip code, bedrooms, and bathrooms have helped increase the sales price.*

```
summary(sale_price_lm)
```

```
##
## Call:
## lm(formula = sale_price ~ zip_code + square_feet_total_living +
##     bedrooms + bath_full_count, data = Housing_df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1759242  -117826   -41208    43890   3832749
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -1.999e+08  1.854e+08  -1.078    0.281
## zip_code        2.041e+03  1.891e+03    1.079    0.281
## square_feet_total_living  1.837e+02  4.377e+00  41.964 < 2e-16 ***
## bedrooms      -2.466e+04  4.446e+03  -5.547 2.96e-08 ***
## bath_full_count  4.195e+04  5.695e+03   7.366 1.87e-13 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 359000 on 12860 degrees of freedom
## Multiple R-squared:  0.2121, Adjusted R-squared:  0.2118
## F-statistic: 865.3 on 4 and 12860 DF,  p-value: < 2.2e-16
```

**Considering the parameters of the multiple regression model you have created. What are the standardized betas for each parameter and what do the values indicate?**

```
install.packages("car")
```

```
##
## The downloaded binary packages are in
## /var/folders/wm/_x82v16s45bgfptwshv23l3r0000gp/T//RtmpNkjeiI/downloaded_packages
```

```
library(car)
```

*## I used the comparecoefs to dive deeper into this question and found that the the most significant difference lied within the bedrooms and full baths.*

```
compareCoefs(square_foot_price_lm, sale_price_lm)
```

```
## Calls:
## 1: lm(formula = sale_price ~ sq_ft_lot, data = Housing_df)
## 2: lm(formula = sale_price ~ zip_code + square_feet_total_living + bedrooms
##     + bath_full_count, data = Housing_df)
##
```

```
##               Model 1   Model 2
## (Intercept)   6.42e+05 -2.00e+08
## SE            3.80e+03  1.85e+08
##
## sq_ft_lot      0.8510
## SE            0.0622
##
## zip_code              2041
## SE                   1891
##
## square_feet_total_living 183.66
## SE                   4.38
##
## bedrooms          -24663
## SE                4446
##
## bath_full_count    41949
## SE                5695
##
```

Calculate the confidence intervals for the parameters in your model and explain what the results indicate.

```
library(MASS, pos = 15)
## This shows us a 95% confidence intervals of 631569.6 645448.7 for the parameter sq_ft_lot.
with(Housing_df, (t.test(sale_price, sq_ft_lot, alternative = 'two.sided', conf.level = .95, paired = TRUE)))

##
## Paired t-test
##
## data: sale_price and sq_ft_lot
## t = 180.35, df = 12864, p-value < 2.2e-16
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  631569.6 645448.7
## sample estimates:
## mean of the differences
##                638509.2
```

Assess the improvement of the new model compared to your original model (simple regression model) by testing whether this change is significant by performing an analysis of variance.

```
## I used the compareCoefs again to see if the change is significant and we found that we have a significant difference between model 1 and 2.
compareCoefs(square_foot_price_lm, sale_price_lm)

## Calls:
## 1: lm(formula = sale_price ~ sq_ft_lot, data = Housing_df)
```

```
## 2: lm(formula = sale_price ~ zip_code + square_feet_total_living + bedroom
s
##   + bath_full_count, data = Housing_df)
##
##               Model 1    Model 2
## (Intercept)    6.42e+05 -2.00e+08
## SE              3.80e+03  1.85e+08
##
## sq_ft_lot      0.8510
## SE              0.0622
##
## zip_code                2041
## SE                    1891
##
## square_feet_total_living    183.66
## SE                        4.38
##
## bedrooms                -24663
## SE                    4446
##
## bath_full_count           41949
## SE                    5695
##
```

*## I also created anovas for each model separately to see the differences on their own.*

*## In this test we will be performing an analysis of variance on the square\_foot\_price\_lm*

```
square_footAnova <- aov(sale_price ~ sq_ft_lot, data = Housing_df)
```

*## We see that the Estimated effects may be unbalanced as the Residual standard error: 401483.8 which shows a significant room for error*

```
square_footAnova
```

```
## Call:
```

```
##   aov(formula = sale_price ~ sq_ft_lot, data = Housing_df)
```

```
##
```

```
## Terms:
```

```
##               sq_ft_lot    Residuals
```

```
## Sum of Squares  3.019674e+13 2.073377e+15
```

```
## Deg. of Freedom      1      12863
```

```
##
```

```
## Residual standard error: 401483.8
```

```
## Estimated effects may be unbalanced
```

*## In this test we will be performing an analysis of variance on the sale\_price\_lm*

```
SaleAnova <- aov(sale_price ~ zip_code + square_feet_total_living + bedrooms
+ bath_full_count, data = Housing_df)
```

*## We see that the Estimated effects may be unbalanced as the Residual standard error: 401483.8 which shows a significant room for error*

*rd error: 359008.6*

SaleAnova

```
## Call:
## aov(formula = sale_price ~ zip_code + square_feet_total_living +
##     bedrooms + bath_full_count, data = Housing_df)
##
## Terms:
##               zip_code square_feet_total_living      bedrooms
## Sum of Squares 7.610438e+12      4.276822e+14 3.798816e+12
## Deg. of Freedom      1              1              1
##               bath_full_count      Residuals
## Sum of Squares 6.992478e+12 1.657490e+15
## Deg. of Freedom      1      12860
##
## Residual standard error: 359008.6
## Estimated effects may be unbalanced
```

**Perform casewise diagnostics to identify outliers and/or influential cases, storing each function's output in a dataframe assigned to a unique variable name.**

*## I created a data frame that is assigned to a unique variable name known as uniquehousing which we see the Coefficients for each variable we could enter.*

```
install.packages("outliers")
```

```
##
```

```
## The downloaded binary packages are in
```

```
## /var/folders/wm/_x82v16s45bgfptwshv23l3r0000gp/T//RtmpNkjeiI/downloaded_packages
```

```
library(outliers)
```

```
uniquehousing <- lm(formula = sale_price ~ sale_reason + sale_instrument + zip_code + building_grade + square_feet_total_living + bedrooms + bath_full_count + bath_half_count + bath_3qtr_count + year_built + year_renovated + sq_ft_lot, data = Housing_df)
```

```
uniquehousing
```

```
##
```

```
## Call:
```

```
## lm(formula = sale_price ~ sale_reason + sale_instrument + zip_code +
##     building_grade + square_feet_total_living + bedrooms + bath_full_count +
##     bath_half_count + bath_3qtr_count + year_built + year_renovated +
##     sq_ft_lot, data = Housing_df)
##
```

```
## Coefficients:
```

```
##               (Intercept)                sale_reason                sale_instrument
```

```
##          -3.931e+07          -1.179e+04          2.430e+
02
##          zip_code          building_grade  square_feet_total_livi
ng
##          3.507e+02          3.077e+04          1.443e+
02
##          bedrooms          bath_full_count          bath_half_cou
nt
##          -3.718e+03          -6.489e+01          -9.841e+
02
##          bath_3qtr_count          year_built          year_renovat
ed
##          -1.568e+04          2.507e+03          7.629e+
01
##          sq_ft_lot
##          2.927e-01
```

*## Next I perform an outlier test on the new data frame and the other two existing ones.*

```
outlierTest(uniquehousing)
```

```
##          rstudent unadjusted p-value Bonferroni p
## 11992 10.57364          5.0505e-26  6.4975e-22
## 6430 10.46158          1.6418e-25  2.1122e-21
## 6438 10.44803          1.8918e-25  2.4339e-21
## 6437 10.42057          2.5199e-25  3.2418e-21
## 4649 10.40985          2.8176e-25  3.6248e-21
## 6431 10.32263          6.9642e-25  8.9594e-21
## 6436 10.29156          9.5952e-25  1.2344e-20
## 6441 10.28152          1.0640e-24  1.3689e-20
## 6432 10.25505          1.3969e-24  1.7971e-20
## 6442 10.21695          2.0641e-24  2.6555e-20
```

*## Next I will remove outliers as seen below:*

```
Updated_Housing_df <- Housing_df[-c(11992,6430,6438,6437,4649,6431,6436,6441,
6432,6442),]
```

*## Next I perform an outlier test on the square\_foot\_price\_lm*

```
outlierTest(square_foot_price_lm)
```

```
##          rstudent unadjusted p-value Bonferroni p
## 6438 9.334760          1.1763e-20  1.5134e-16
## 6437 9.334494          1.1793e-20  1.5171e-16
## 6441 9.334316          1.1813e-20  1.5197e-16
## 6433 9.334031          1.1844e-20  1.5237e-16
## 6434 9.333823          1.1867e-20  1.5267e-16
## 6430 9.333677          1.1884e-20  1.5288e-16
```



```
## 6442 9.332473      1.2018e-20    1.5462e-16
## 6439 9.331469      1.2132e-20    1.5608e-16
## 6431 9.331388      1.2141e-20    1.5620e-16
## 6429 9.329466      1.2362e-20    1.5904e-16
```

*## Next I will remove outliers as seen below:*

```
Updated_Housing_df <- Housing_df[-c(6438,6437,6441,6433,6434,6430,6442,6439,6431,6429),]
```

*## Next I perform an outlier test on the sale\_price\_lm*

```
outlierTest(sale_price_lm)
```

```
##          rstudent unadjusted p-value Bonferroni p
## 11992 10.72424      1.0160e-26    1.3070e-22
## 4649  10.59209      4.1537e-26    5.3438e-22
## 6430  10.46710      1.5493e-25    1.9932e-21
## 6437  10.41523      2.6637e-25    3.4269e-21
## 6438  10.40265      3.0366e-25    3.9066e-21
## 6431  10.31155      7.8074e-25    1.0044e-20
## 6436  10.28046      1.0756e-24    1.3838e-20
## 6432  10.17689      3.1066e-24    3.9967e-20
## 6433  10.14913      4.1211e-24    5.3018e-20
## 6434  10.14913      4.1211e-24    5.3018e-20
```

*## Next I will remove outliers as seen below:*

```
Updated_Housing_df <- Housing_df[-c(11992,4649,6430,6437,6438,6431,6436,6432,6433,6434),]
```

*## I shall show the updated Updated\_Housing\_df*

```
str(Updated_Housing_df)
```

```
## tibble[,24] [12,855 × 24] (S3: tbl_df/tbl/data.frame)
## $ sale_date      : chr [1:12855] "1/3/06" "1/3/06" "1/3/06" "1/3/06" ...
## $ sale_price     : num [1:12855] 698000 649990 572500 420000 369900 ...
## $ sale_reason    : num [1:12855] 1 1 1 1 1 1 1 1 1 1 ...
## $ sale_instrument : num [1:12855] 3 3 3 3 3 15 3 3 3 3 ...
## $ sale_warning   : chr [1:12855] "No" "No" "No" "No" ...
## $ sitetype       : chr [1:12855] "R1" "R1" "R1" "R1" ...
## $ addr_full      : chr [1:12855] "17021 NE 113TH CT" "11927 178TH PL NE" "13315 174TH AVE NE" "3303 178TH AVE NE" ...
## $ zip_code       : num [1:12855] 98052 98052 98052 98052 98052 .
## ..
## $ ctyname        : chr [1:12855] "REDMOND" "REDMOND" "REDMOND" "REDMOND" ...
## $ postalctyn     : chr [1:12855] "REDMOND" "REDMOND" "REDMOND" "REDMOND" "
```

```

REDMOND" ...
## $ lon : num [1:12855] -122 -122 -122 -122 -122 ...
## $ lat : num [1:12855] 47.7 47.7 47.7 47.6 47.7 ...
## $ building_grade : num [1:12855] 9 9 8 8 7 7 10 10 9 8 ...
## $ square_feet_total_living: num [1:12855] 2810 2880 2770 1620 1440 4160 3
960 3720 4160 2760 ...
## $ bedrooms : num [1:12855] 4 4 4 3 3 4 5 4 4 4 ...
## $ bath_full_count : num [1:12855] 2 2 1 1 1 2 3 2 2 1 ...
## $ bath_half_count : num [1:12855] 1 0 1 0 0 1 0 1 1 0 ...
## $ bath_3qtr_count : num [1:12855] 0 1 1 1 1 1 1 0 1 1 ...
## $ year_built : num [1:12855] 2003 2006 1987 1968 1980 ...
## $ year_renovated : num [1:12855] 0 0 0 0 0 0 0 0 0 0 ...
## $ current_zoning : chr [1:12855] "R4" "R4" "R6" "R4" ...
## $ sq_ft_lot : num [1:12855] 6635 5570 8444 9600 7526 ...
## $ prop_type : chr [1:12855] "R" "R" "R" "R" ...
## $ present_use : num [1:12855] 2 2 2 2 2 2 2 2 2 2 ...

```

*## updated uniquehousing data frame*

```

updated_square_foot_price_lm <- lm(sale_price ~ sq_ft_lot, data = Updated_Housing_df)

```

*## After that I will updated the lm with all the variables I used*

```

reupdated_square_foot_price_lm <- lm(sale_price ~ sale_reason + sale_instrument + zip_code + building_grade + square_feet_total_living + bedrooms + bath_full_count + bath_half_count + bath_3qtr_count + year_built + year_renovated + sq_ft_lot, data = Updated_Housing_df)

```

```

reupdated_square_foot_price_lm

```

```

##
## Call:
## lm(formula = sale_price ~ sale_reason + sale_instrument + zip_code +
##      building_grade + square_feet_total_living + bedrooms + bath_full_count +
##      bath_half_count + bath_3qtr_count + year_built + year_renovated +
##      sq_ft_lot, data = Updated_Housing_df)
##
## Coefficients:
##              (Intercept)              sale_reason              sale_instrument
##              -1.126e+08              -9.407e+03              -4.161e+03
##              zip_code              building_grade square_feet_total_living
##              1.104e+03              3.276e+04              1.466e+02
##              bedrooms              bath_full_count              bath_half_count

```

##	-8.297e+03	4.313e+03	-1.863e+
03			
##	bath_3qtr_count	year_built	year_renovat
ed			
##	-1.168e+04	2.208e+03	7.451e+
01			
##	sq_ft_lot		
##	1.484e-01		

Calculate the standardized residuals using the appropriate command, specifying those that are +2, storing the results of large residuals in a variable you create.

*## I decided to update the hatvalues, hatvalues, covariance.ratios, standardized.residuals, studentized.residuals, cooks.distance, and dfbeta.*

```
Updated_Housing_df$dfbeta <- dfbeta(reupdated_square_foot_price_lm)
Updated_Housing_df$leverage <- hatvalues(reupdated_square_foot_price_lm)
Updated_Housing_df$covariance.ratios <- covratio(reupdated_square_foot_price_lm)
Updated_Housing_df$standardized.residuals <- rstandard(reupdated_square_foot_price_lm)
Updated_Housing_df$studentized.residuals <- rstudent(reupdated_square_foot_price_lm)
Updated_Housing_df$cooks.distance <- cooks.distance(reupdated_square_foot_price_lm)
```

*## I used the str() function to calculate the standardized residuals*

```
str(Updated_Housing_df)
```

```
## tibble[,30] [12,855 × 30] (S3: tbl_df/tbl/data.frame)
## $ sale_date          : chr [1:12855] "1/3/06" "1/3/06" "1/3/06" "1/3/06" ...
## $ sale_price         : num [1:12855] 698000 649990 572500 420000 369900 ...
## $ sale_reason        : num [1:12855] 1 1 1 1 1 1 1 1 1 1 ...
## $ sale_instrument    : num [1:12855] 3 3 3 3 3 15 3 3 3 3 ...
## $ sale_warning       : chr [1:12855] "No" "No" "No" "No" ...
## $ sitetype           : chr [1:12855] "R1" "R1" "R1" "R1" ...
## $ addr_full          : chr [1:12855] "17021 NE 113TH CT" "11927 178TH PL NE" "13315 174TH AVE NE" "3303 178TH AVE NE" ...
## $ zip_code           : num [1:12855] 98052 98052 98052 98052 98052 .
## ..
## $ ctyname            : chr [1:12855] "REDMOND" "REDMOND" "REDMOND" "REDMOND" ...
## $ postalctyn         : chr [1:12855] "REDMOND" "REDMOND" "REDMOND" "REDMOND" ...
## $ lon                : num [1:12855] -122 -122 -122 -122 -122 ...
## $ lat                : num [1:12855] 47.7 47.7 47.7 47.6 47.7 ...
## $ building_grade     : num [1:12855] 9 9 8 8 7 7 10 10 9 8 ...
## $ square_feet_total_living: num [1:12855] 2810 2880 2770 1620 1440 4160 3960 3720 4160 2760 ...
## $ bedrooms           : num [1:12855] 4 4 4 3 3 4 5 4 4 4 ...
```

```
## $ bath_full_count      : num [1:12855] 2 2 1 1 1 2 3 2 2 1 ...
## $ bath_half_count      : num [1:12855] 1 0 1 0 0 1 0 1 1 0 ...
## $ bath_3qtr_count      : num [1:12855] 0 1 1 1 1 1 1 0 1 1 ...
## $ year_built           : num [1:12855] 2003 2006 1987 1968 1980 ...
## $ year_renovated       : num [1:12855] 0 0 0 0 0 0 0 0 0 0 ...
## $ current_zoning       : chr [1:12855] "R4" "R4" "R6" "R4" ...
## $ sq_ft_lot            : num [1:12855] 6635 5570 8444 9600 7526 ...
## $ prop_type            : chr [1:12855] "R" "R" "R" "R" ...
## $ present_use          : num [1:12855] 2 2 2 2 2 2 2 2 2 2 ...
## $ dfbeta               : num [1:12855, 1:13] -88130 -167917 -61105 -25
431 -32783 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : chr [1:12855] "1" "2" "3" "4" ...
## .. ..$ : chr [1:13] "(Intercept)" "sale_reason" "sale_instrument" "zip_c
ode" ...
## $ leverage             : Named num [1:12855] 0.000275 0.00037 0.000382
0.000451 0.000304 ...
## ..- attr(*, "names")= chr [1:12855] "1" "2" "3" "4" ...
## $ covariance.ratios    : Named num [1:12855] 1 1 1 1 1 ...
## ..- attr(*, "names")= chr [1:12855] "1" "2" "3" "4" ...
## $ standardized.residuals : Named num [1:12855] -0.151 -0.313 -0.256 -0.1
16 -0.166 ...
## ..- attr(*, "names")= chr [1:12855] "1" "2" "3" "4" ...
## $ studentized.residuals : Named num [1:12855] -0.151 -0.312 -0.256 -0.1
16 -0.166 ...
## ..- attr(*, "names")= chr [1:12855] "1" "2" "3" "4" ...
## $ cooks.distance       : Named num [1:12855] 4.81e-07 2.78e-06 1.93e-0
6 4.65e-07 6.45e-07 ...
## ..- attr(*, "names")= chr [1:12855] "1" "2" "3" "4" ...
```

### Use the appropriate function to show the sum of large residuals.

*## I used large.residual, standardized.residuals > 2, and studentized.residuals < -2 to show the sum of large residuals*

```
Updated_Housing_df$large.residual <- Updated_Housing_df$standardized.residual
s > 2 | Updated_Housing_df$studentized.residuals < -2
str(Updated_Housing_df)
```

```
## tibble[,31] [12,855 × 31] (S3: tbl_df/tbl/data.frame)
## $ sale_date           : chr [1:12855] "1/3/06" "1/3/06" "1/3/06" "1/3
/06" ...
## $ sale_price          : num [1:12855] 698000 649990 572500 420000 369
900 ...
## $ sale_reason         : num [1:12855] 1 1 1 1 1 1 1 1 1 1 ...
## $ sale_instrument     : num [1:12855] 3 3 3 3 3 15 3 3 3 3 ...
## $ sale_warning        : chr [1:12855] "No" "No" "No" "No" ...
## $ sitetype            : chr [1:12855] "R1" "R1" "R1" "R1" ...
## $ addr_full           : chr [1:12855] "17021 NE 113TH CT" "11927 178T
H PL NE" "13315 174TH AVE NE" "3303 178TH AVE NE" ...
## $ zip_code            : num [1:12855] 98052 98052 98052 98052 98052 .
..
```

```

## $ ctyname : chr [1:12855] "REDMOND" "REDMOND" "REDMOND" "
REDMOND" ...
## $ postalctyn : chr [1:12855] "REDMOND" "REDMOND" "REDMOND" "
REDMOND" ...
## $ lon : num [1:12855] -122 -122 -122 -122 -122 ...
## $ lat : num [1:12855] 47.7 47.7 47.7 47.6 47.7 ...
## $ building_grade : num [1:12855] 9 9 8 8 7 7 10 10 9 8 ...
## $ square_feet_total_living: num [1:12855] 2810 2880 2770 1620 1440 4160 3
960 3720 4160 2760 ...
## $ bedrooms : num [1:12855] 4 4 4 3 3 4 5 4 4 4 ...
## $ bath_full_count : num [1:12855] 2 2 1 1 1 2 3 2 2 1 ...
## $ bath_half_count : num [1:12855] 1 0 1 0 0 1 0 1 1 0 ...
## $ bath_3qtr_count : num [1:12855] 0 1 1 1 1 1 1 0 1 1 ...
## $ year_built : num [1:12855] 2003 2006 1987 1968 1980 ...
## $ year_renovated : num [1:12855] 0 0 0 0 0 0 0 0 0 0 ...
## $ current_zoning : chr [1:12855] "R4" "R4" "R6" "R4" ...
## $ sq_ft_lot : num [1:12855] 6635 5570 8444 9600 7526 ...
## $ prop_type : chr [1:12855] "R" "R" "R" "R" ...
## $ present_use : num [1:12855] 2 2 2 2 2 2 2 2 2 2 ...
## $ dfbeta : num [1:12855, 1:13] -88130 -167917 -61105 -25
431 -32783 ...
## ... attr(*, "dimnames")=List of 2
## .. .$ : chr [1:12855] "1" "2" "3" "4" ...
## .. .$ : chr [1:13] "(Intercept)" "sale_reason" "sale_instrument" "zip_c
ode" ...
## $ leverage : Named num [1:12855] 0.000275 0.00037 0.000382
0.000451 0.000304 ...
## ... attr(*, "names")= chr [1:12855] "1" "2" "3" "4" ...
## $ covariance.ratios : Named num [1:12855] 1 1 1 1 1 ...
## ... attr(*, "names")= chr [1:12855] "1" "2" "3" "4" ...
## $ standardized.residuals : Named num [1:12855] -0.151 -0.313 -0.256 -0.1
16 -0.166 ...
## ... attr(*, "names")= chr [1:12855] "1" "2" "3" "4" ...
## $ studentized.residuals : Named num [1:12855] -0.151 -0.312 -0.256 -0.1
16 -0.166 ...
## ... attr(*, "names")= chr [1:12855] "1" "2" "3" "4" ...
## $ cooks.distance : Named num [1:12855] 4.81e-07 2.78e-06 1.93e-0
6 4.65e-07 6.45e-07 ...
## ... attr(*, "names")= chr [1:12855] "1" "2" "3" "4" ...
## $ large.residual : Named logi [1:12855] FALSE FALSE FALSE FALSE
FALSE FALSE ...
## ... attr(*, "names")= chr [1:12855] "1" "2" "3" "4" ...

```

**Which specific variables have large residuals (only cases that evaluate as TRUE)?**

*## Next we will use the sum to show the sum of the large residuals that evaluate as TRUE to be 330*

```
sum(Updated_Housing_df$large.residual)
```

```
## [1] 330
```

*## After we have received our specific variables we can put it into a table.*

```
Updated_Housing_df[Updated_Housing_df$large.residual , c("sale_price", "sale_reason", "sale_instrument", "zip_code", "building_grade", "square_feet_total_living", "bedrooms", "bath_full_count", "bath_half_count", "bath_3qtr_count", "year_built", "year_renovated", "sq_ft_lot")]
```

```
## # A tibble: 330 x 13
```

```
##   sale_price sale_reason sale_instrument zip_code building_grade
##   <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1    265000         1         3      98053         10
## 2   1390000         1         3      98053          6
## 3    390000         1         3      98052         11
## 4   1588359         1         3      98053          9
## 5   1450000         1         3      98052          8
## 6   1450000         1         3      98052          6
## 7    163000         1         3      98053          9
## 8    270000         1         3      98053         11
## 9    200000         1         3      98053         10
## 10   300000         1         3      98052         11
```

```
## # ... with 320 more rows, and 8 more variables: square_feet_total_living <dbl>,
```

```
## #   bedrooms <dbl>, bath_full_count <dbl>, bath_half_count <dbl>,
```

```
## #   bath_3qtr_count <dbl>, year_built <dbl>, year_renovated <dbl>,
```

```
## #   sq_ft_lot <dbl>
```

### Investigate further by calculating the leverage, cooks distance, and covariance ratios. Comment on all cases that are problematic

*## Next I use the Updated\_Housing\_df to calculate the Leverage, cooks distance, and covariance ratios*

```
Updated_Housing_df[Updated_Housing_df$large.residual , c("leverage" , "cooks.distance", "covariance.ratios")]
```

```
## # A tibble: 330 x 3
```

```
##   leverage cooks.distance covariance.ratios
##   <dbl>      <dbl>      <dbl>
## 1 0.00135    0.000669    0.996
## 2 0.00296    0.00238    0.993
## 3 0.00142    0.000738    0.996
## 4 0.000787   0.000285    0.997
## 5 0.00101    0.000324    0.998
## 6 0.00233    0.00245    0.990
## 7 0.00119    0.000601    0.996
## 8 0.183      0.186      1.21
## 9 0.00535    0.00497    0.994
## 10 0.000707   0.000278    0.997
```

```
## # ... with 320 more rows
```

Perform the necessary calculations to assess the assumption of independence and state if the condition is met or not.

*## I used the durbin Watson test DWT to determine if the necessary calculations to assess the assumption of independence and state if the condition is met and I conclude yes it is possible as the Durbin Watson test is about the correlation of the residuals. If the data are ordered in some way, you'll get a significant DW test*

```
dwt(reupdated_square_foot_price_lm)

## lag Autocorrelation D-W Statistic p-value
## 1 0.7215049 0.5569814 0
## Alternative hypothesis: rho != 0
```

Perform the necessary calculations to assess the assumption of no multicollinearity and state if the condition is met or not.

*## In this calculation we can assess multicollinearity by computing the variance inflation factor known as VIF which measures how much the variance of a regression coefficient is inflated due to multicollinearity in the model. As the smallest possible value of VIF is one which is the absence of multicollinearity. While a VIF value that exceeds 5 or 10 indicates a problematic amount of collinearity.*

```
car::vif(reupdated_square_foot_price_lm)

##          sale_reason          sale_instrument          zip_code
##          1.211430          1.220892          1.039035
##    building_grade square_feet_total_living          bedrooms
##          2.416949          4.223336          1.756799
##    bath_full_count    bath_half_count    bath_3qtr_count
##          2.509261          1.440674          2.084983
##          year_built    year_renovated    sq_ft_lot
##          1.587690          1.089389          1.174761
```

*## As seen above the VIF score for the predictor variable square\_feet\_total\_living is the highest we see with a VIF = 4.22 which would be the only one that is close enough to 5 that could even be thought to might be problematic. But in this case I believe none of the variables in my lm would be problematic.*

Visually check the assumptions related to the residuals using the plot() and hist() functions. Summarize what each graph is informing you of and if any anomalies are present.

*## In this graph we see the arg frequency that is related to the residual of the updated housing data frame*

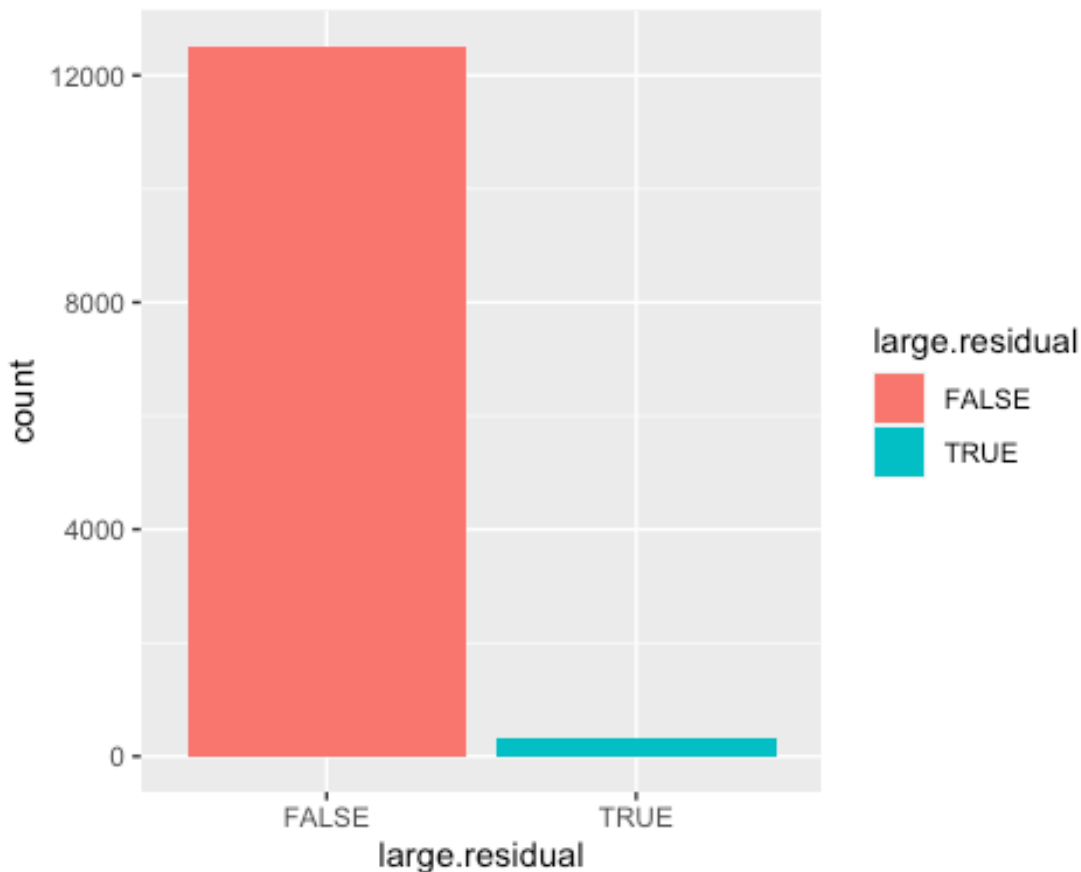
```
install.packages("Rcmdr")
```

```
##
## The downloaded binary packages are in
```

```
## /var/folders/wm/_x82v16s45bgfptwshv23l3r0000gp/T//RtmpNkjeiI/downloaded_packages
```

*## In this bar chart we saw that out of all our properties we found that over 12,000 are not considered large residual with a small amount being large residual.*

```
ggplot(Updated_Housing_df, aes(large.residual))+geom_bar(aes(fill = large.residual))
```



```
with(Updated_Housing_df, Hist(standardized.residuals, scale="frequency", col="blue", xlab="Residuals"))
```

```
## Error in Hist(standardized.residuals, scale = "frequency", col = "blue", :  
could not find function "Hist"
```



Overall, is this regression model unbiased? If an unbiased regression model, what does this tell us about the sample vs. the entire population model?

I believe the regression model is unbiased as we see in our VIF that all our variables are non problematic. With this being said in regards to our sample vs the entire population model I would say it can be biased as we can change our sample to only view variables that we would know that would pair well together as the entire population we would see outliers and data that doesn't really reflect our needs.